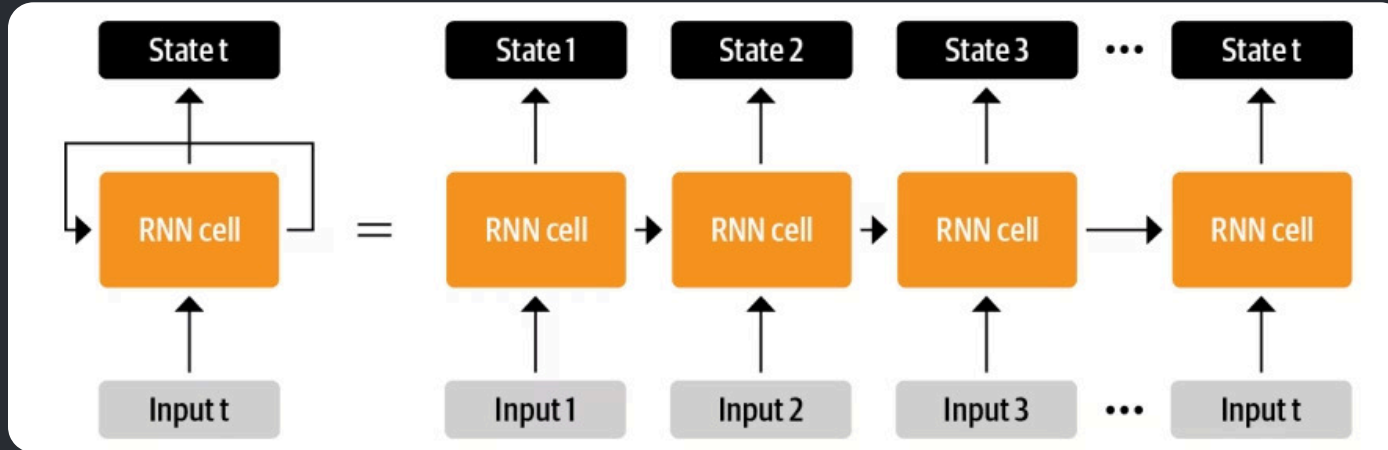


Generative AI: An Overview

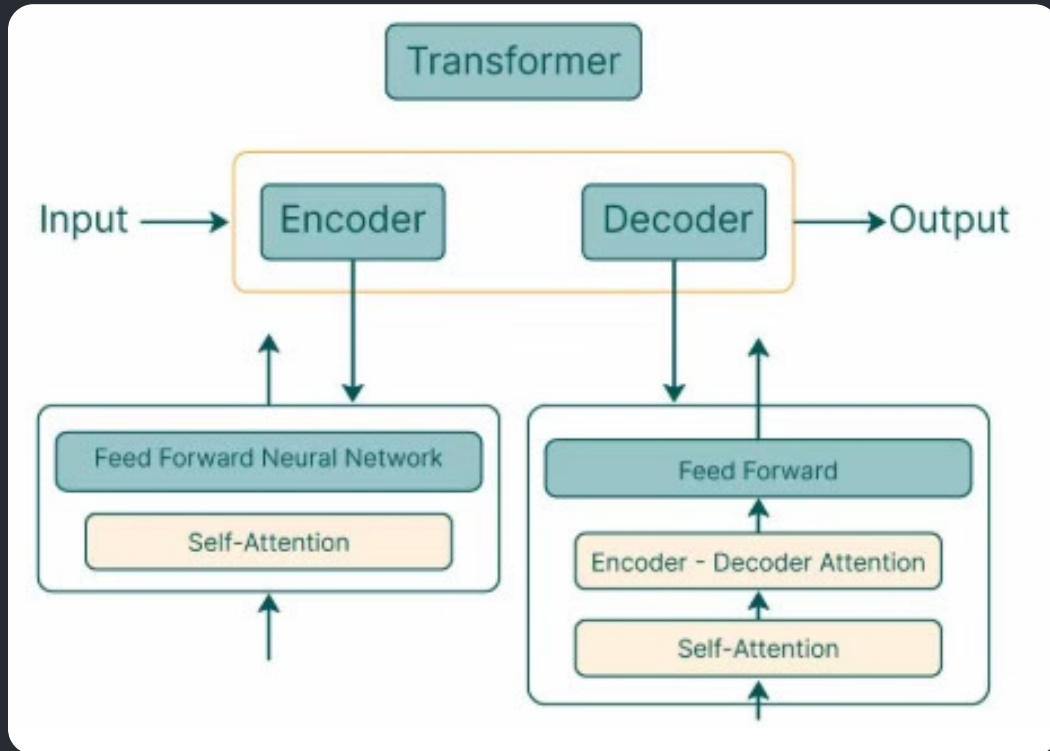


Understanding Recurrent Neural Networks (RNNs)




- RNNs are a type of neural network. They are designed to process sequential data.
- These architectures were widely used for NLP tasks, speech processing, and time series.
- Challenge-?

The Rise of Transformers: Self-Attention



- In 2017, researchers at Google published a paper that proposed a novel neural network architecture for sequence modeling known as Transformer.
- Outperformed recurrent neural networks (RNNs) on machine translation tasks, both in terms of translation quality and training cost.



A Timeline of Large Language Models

2022: ChatGPT

Generative Pre-trained Transformer 2.

2024: Meta's Llama 3, Claude 3, and Q2, and Mistral's Mixtral 8x7B

Larger and more powerful model.

2025: DeepSeek-R1

Multimodality: Text, Image, Video

Diving into ChatGPT

Generative

Next word prediction

Pre-trained

LLM is pre-trained on massive amount of text

Transformer

Encoder-decoder architecture

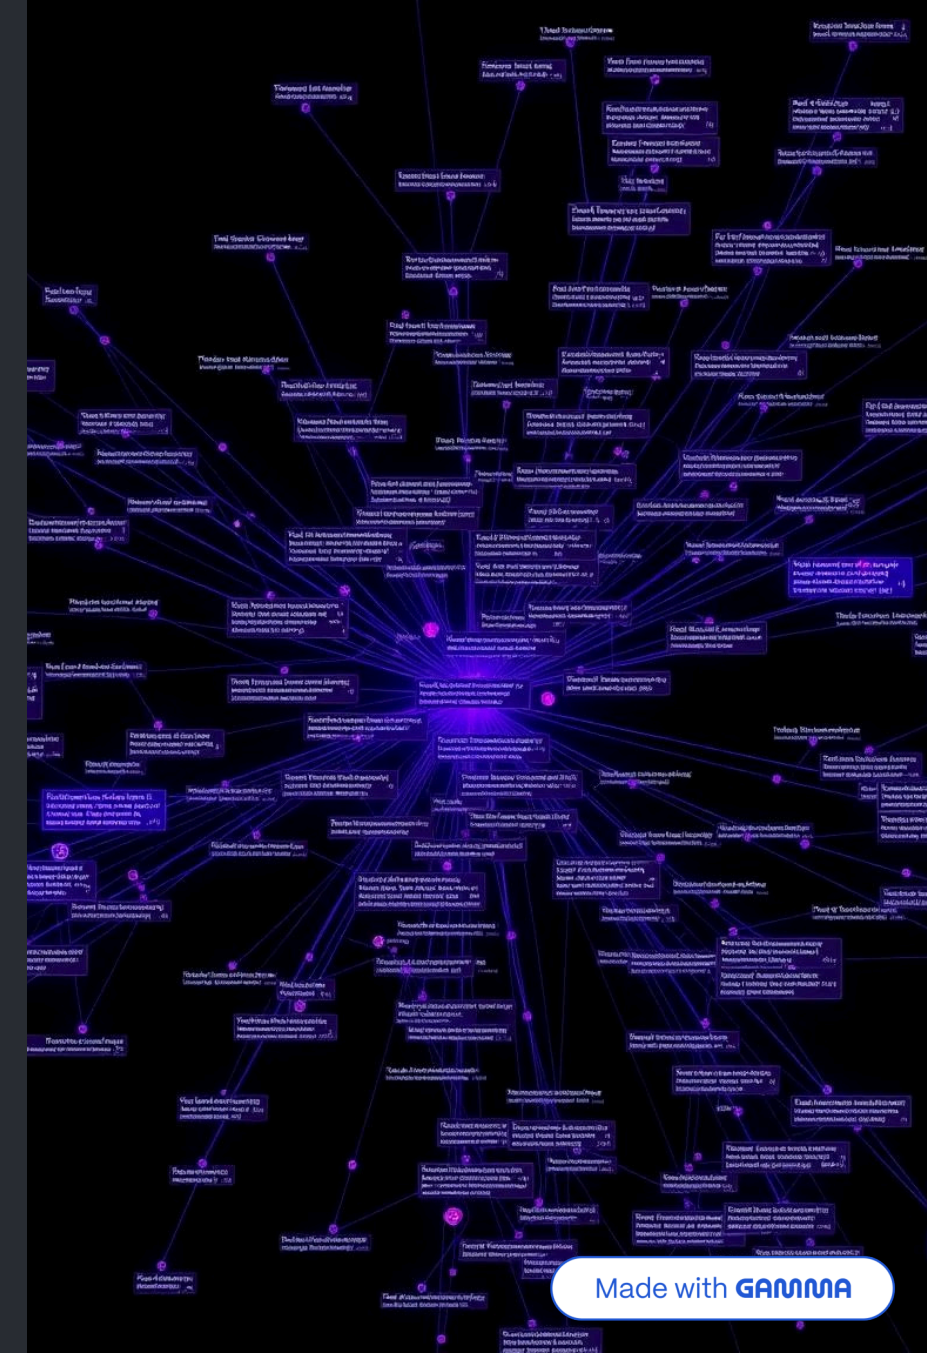
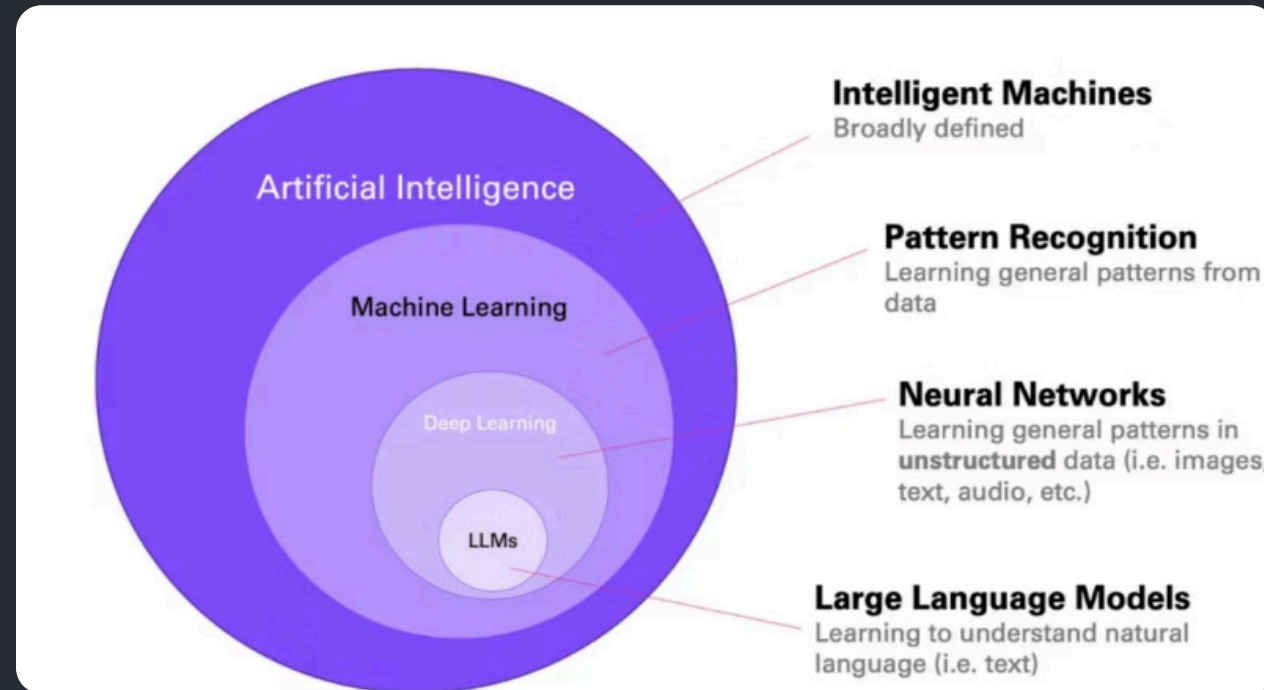
Why did ChatGPT couldn't replace Google Search?

How was ChatGPT trained?



Large Language Models

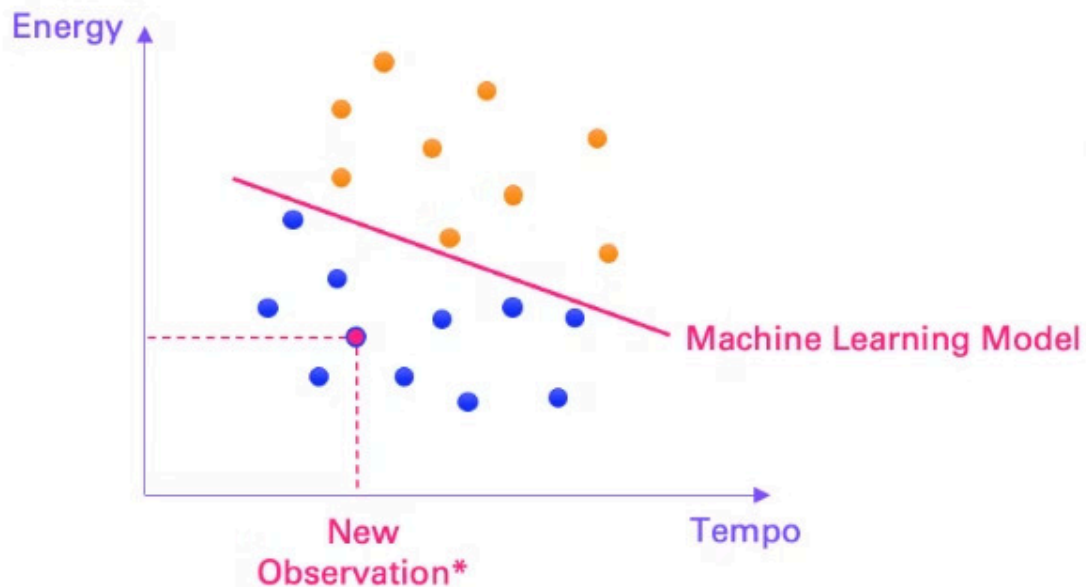
What do LLMs essentially do?



LLMs as Machine Learning Task?

Classification Example: Predicting Music Genre

● R&B Songs ● Reggaeton Songs ● New Observation



*New Observation:

Tempo: 30, Energy: 20

ML

Class	Probability
R&B	0.92
Reggaeton	0.08


LLMs as Deep Learning Task?

Imagine the following task: Predict the next word in a sequence

[The cat likes to sleep in the _____] → What **word** comes next?

Can we frame this as a ML problem? Yes, it's a **classification** task.

Now we have (say)
~50,000 **classes** (i.e.
words)



Language modeling is learning to predict the next word.

Training Data for LLMs

We can create **vast amounts of sequences** for training a language model

● Context ● Next Word ● Ignored

[The cat likes to sleep in the]
[The cat likes to sleep in the]
[The cat likes to sleep in the]
[The cat likes to sleep in the]
[The cat likes to sleep in the]

We do the same with much longer sequences. For example:

A language model is a probability distribution over sequences of words. [...] Given any sequence of words, the model predicts the **next** ...

Or also with code:

```
def square(number):  
    """Calculates the square of a number."""  
    return number ** 2
```

And as a result - the model becomes incredibly good at **predicting the next word** in any sequence.


Massive amounts of training data can be created relatively easily.

Next word Generation

After training: We can **generate text** by predicting **one word at a time**

A trained language model can

Input



LLM

Word	Probability
speak	0.065
generate	0.072
politics	0.001
...	...
walk	0.003

Output at step 1

Word	Probability
ability	0.002
text	0.084
coherent	0.085
...	...
ideas	0.041

Output at step 2

LLMs are an example of what's called "Generative AI"

Phases of LLM Training

Pre-training

- Massive amount of text data from internet - books, research papers, websites
- Model learns to predict the next word

Instruction fine tuning

- Curating Q n A dataset to train the model to answer questions or instructions
- Model learns to become a helpful assistant

Reinforcement Learning from Human Feedback (RLHF)

- Align the output closer to human like responses
- Responses are updated considering human feedback and preference.

Limitation of LLMs

1. Hallucination
2. Mathematical Problem solving
3. Context window
4. Cost

How to make LLMs respond better?

Zero-Shot

- Give some instructions to solve a task.

Few-Shot

- Give some examples of how to solve a task.

Chain-of-Thought(CoT)

- For complex tasks- prompt an LLM to “think step by step”

Latest LLMs & Frameworks

LLMs

Mistral

Mixtral

Llama

Gemini

DeepSeek

Frameworks

Together AI- <https://www.together.ai/>

Groq- <https://groq.com/>

Replicate- <https://replicate.com/>

LiteLLM - <https://www.litellm.ai/>

Hugging Face- <https://huggingface.co/>

Generative AI Project Lifecycle

